

## Segmented Parallel Scan Algorithm

Template Project Name: SegmentedParallelScan

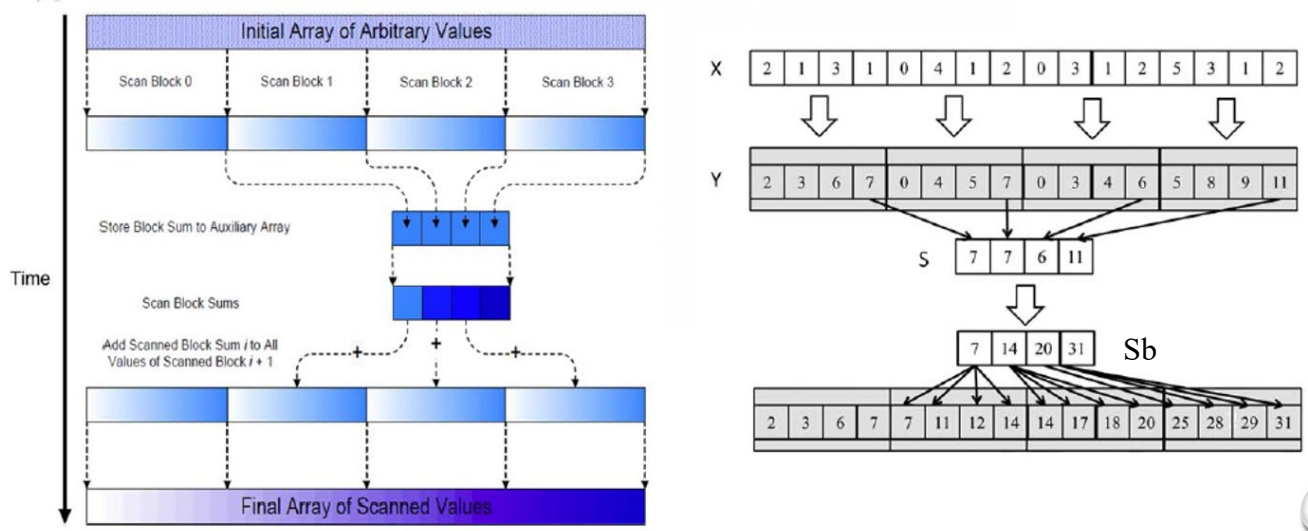


Figure 1 Segmented Parallel Scan Architecture

### Information:

- Predefined constant **B** depicts the block size. The block sizes can be 256 or 512 or 1024
- Predefined constant **VECTOR\_SIZE** is the product of **B** and **Z**, where **Z** is an integer multiplier
- The value of the **Z** has to be less than or equal to **B**
- `__global__ void SingleBlockScan(float* In, float* Out, const int SIZE)`  
GPU Kernel to perform a parallel scan of only a single block with a maximum of 1024 elements using the Efficient Kogge Stone Algorithm. It is the same kernel that was implemented during the class and can only scan a single block
- `__global__ void MultipleBlockScanIntermediateOutput(float* In, float* AuxOut, float* blockSum, const int SIZE)`  
GPU Kernel to perform a scan of individual blocks and produce intermediate output and output of the last element in each block. This kernel should have Vector **X** as the input and produce Vectors **Y** and **S** in Figure 1.
- The vector **Sb** in Figure 1 should be generated using the SingleBlockScan kernel
- `__global__ void AddBlockSumValues(float* AuxIn, float* blockSum, const int SIZE)`  
GPU Kernel to add the block sum of the last element to the Intermediate Outputs to produce the final answer. This kernel should have Vectors **Y** and **Sb** as inputs and produce the final output vector in Figure 1 representing the Parallel Scan result.

**Coding Tasks:**

Implement code of the functions listed below:

- `__global__ void MultipleBlockScanIntermediateOutput(float* In, float* AuxOut, float* blockSum, const int SIZE)`
- `__global__ void AddBlockSumValues(float* AuxIn, float* blockSum, const int SIZE)`
- `__host__ void Helper_Scan(float* Input, float* Output, float* RefOutputData, int SIZE)`

**Note: Examine the function *OnInitializeInputData* to help with debugging.**

**Submissions:**

- Project with commented code (**Note: If the code implemented by you is not commented appropriately, a penalty of 10 points is assigned**) **Points: 35**

- Collect the performance data for the following data sizes: **Points: 10**

B	Z	Execution Times				Total GPU (Only Kernel) Time (usecs)	Speedup
		CPU (msecs)	MultipleBlockScan (usecs)	SingleBlockScan (usecs)	AddBlock (usecs)		
256	64						
256	128						
256	256						
512	128						
512	256						
512	512						
1024	256						
1024	512						
1024	1024						

- Explain the accuracy of the GPU computation with increasing vector size and determine the precision (decimal place) matching the results of the CPU computation. **Points: 5**